



Signal-to-Action Agent

An AI-native Enterprise Revenue Operating System

Built by Team VentureOS · Live demo: SMB Growth & Seller Action · v3

TRACK A · AGENTIC WORKFLOWS

Signals

Multi-Agent

NVIDIA Runtime

Decision Ledger

Human Approval

CEO-level thesis

Enterprises do not need another dashboard or chatbot. They need an accelerated, observable, governed operating system that moves safely from signal → decision → human-approved business outcome — now evolving toward voice-native.



Contents

The original application structure — evolved over two weeks into a working, deployed product.

01 · Team introduction

Why this team can execute a business-grade agentic workload.

02 · AI use case

Product workflow, demo, dataset, evaluation, architecture, governance, and execution.

03 · NVIDIA workload & stack

Provider-abstracted today; NVIDIA NIM / Nemotron / NeMo / Triton-ready by configuration.

04 · Roadmap & mentor asks

Three horizons — current, hackathon (voice), and future — plus mentor asks.

Deck intent

This is version 3 of our original submission. The architecture is unchanged in spirit; the product has been built, deployed, and hardened — and is now evolving toward a voice-native Chief of Staff.



TEAM INTRODUCTION

SKILLSET / EXPERIENCE



Team introduction

03

Product strategy + engineering execution — two builders who shipped a deployed product in two weeks.

Amit Pandey · Product / AI Architecture

Principal PM experience across enterprise platforms, agentic AI prototypes, customer signal systems, experimentation, governance, and AI product storytelling. Two-time global hackathon winner and top-builder recognition in India.

Saurabh Talreja · Engineering / Full-stack Build

Engineering partner focused on implementation, frontend/backend integration, agent runtime wiring, API integration, testing, performance tuning, and demo readiness.

Agentic AI skills

- Multi-agent workflow design
- RAG/provenance patterns
- Evaluation and governance
- Business workflow automation

What we shipped

- Deployed web app + API
- FastAPI + Next.js, live
- HubSpot CRM integration
- Provider abstraction + BYOK

NVIDIA readiness

- NeMo Agent Toolkit / NemoClaw
- Nemotron / NIM endpoints
- GPU-backed inference
- Mentor-led optimization

Execution principle: business problem first · bounded autonomy · measurable workload · human approval before action




AI Use Case

Expand on your project



AI use case: the Signal-to-Action workflow

One governed loop — from a raw customer signal to an explainable, human-approved business outcome, then back into learning.



Executive Brief & Portfolio Pulse

What changed overnight, why it matters, and what to do — ranked across the entire book of accounts.

Evidence Intelligence

Every recommendation cites the underlying signals and notes — no unexplained scores, no black box.

Decision Ledger

Append-only record of evidence, confidence, caveats, approval, and the outcome that followed.

Continuous Learning

Captured outcomes return as fresh signals — the loop compounds and the system gets sharper over time.

Project motivation: why this matters

Enterprises need an action layer, not another passive dashboard or generic chatbot.

TRACK

A

Growth teams often start Monday with thousands of accounts and scattered signals. The hard part is not seeing data — it is deciding where to act, why to act, and how to act safely.

Current pain

Customer context sits across CRM, spend, engagement, product usage, campaign response, support risk, and follow-up history. Human teams spend time stitching context before acting.

Why agentic AI

A controlled agent workflow can assign reasoning tasks to specialized agents, compare evidence, generate action options, and explain every recommendation.

Why sovereign

Enterprise teams need data boundaries, traceability, audit logs, confidence, and human approval before AI-generated recommendations become business action.

Now shipped

Over two weeks we turned this thesis into a working Enterprise Revenue Operating System — deployed, governed, and now evolving toward voice. Everything that follows is built, not hypothetical.

Demo scenario: SMB Growth & Seller Action

A concrete enterprise workflow that reviewers can understand in five minutes.

TRACK

A

1

Morning brief

Seller opens Signal-to-Action; the Executive Attention Brief shows what changed overnight.

2

Portfolio Pulse

Portfolio Pulse highlights the accounts that moved and why they matter right now.

3

Open account

The Workspace opens the top account; Evidence Intelligence explains the why behind the score.

4

Approve

Recommendation reviewed; the seller approves, edits, or rejects with full evidence in view.

5

Execute & capture

Revenue Execution Center orchestrates the action; outcome captured; Decision Ledger updated.

In the live demo

Live on production: 99 accounts, 108 signals, 6 agents, and 10 ranked recommendations.
Every recommendation carries evidence, confidence, caveats, and a hard human-approval gate.
Signal → brief → recommendation → approval → execution → outcome → ledger, end to end.

Dataset for training and inferencing

Synthetic-by-default plus a HubSpot test CRM — safe, portable, and sovereign by design.

TRACK

A

Live dataset

Synthetic SMB accounts plus a connected HubSpot test CRM: account profile, spend trend, engagement, support, product usage, campaign response, segment, and follow-up history.

Current scale

99 accounts · 108 signals · 6 agents · 10 ranked recommendations live on production — with synthetic generators for larger stress tests.

License and safety

No customer-confidential or employer data. Synthetic-by-default; the HubSpot link is a test portal. Safe to share and fully portable.

Inference & provider plan

Model strategy

Inference-only with provider abstraction: deterministic mock today, NVIDIA NIM / Nemotron when access lands. No fine-tuning required for the workflow.

Why not fine-tune initially

The product proves agent workflow quality — routing, reasoning, structured outputs, evidence validation, latency, and human approval. Fine-tuning is a later optimization, not a dependency.

Evaluation: results and how we measure

Transparent by design: automated evals run green, and every decision is reviewable end to end.

TRACK

A

Current status

Built, deployed, and hardened. Architecture, data schema, demo flow, and an automated evaluation harness all run today — live on Vercel + Render.

What we measure

Schema validity, evidence presence, confidence range, governance caveats, default-pending approval, and latency budget — 10 automated checks, currently 10/10 passing.

Evaluation signals

Agent trace quality, hallucination checks, evidence sufficiency, confidence calibration, default-pending approval, and time-to-action.

What the system proves today

- ❖ Every recommendation cites multiple relevant signals from the live dataset.
- ❖ Every recommended action includes evidence, confidence, caveats, and an explicit human-approval step.
- ❖ Sub-second to low-second end-to-end workflow for a full account-prioritization query.
- ❖ Visible agent execution, decision ledger, and provider abstraction — NVIDIA-ready when access lands.



Tech Stack

Strategy for building/improving your project



Technology stack: what the product runs on

A shipped full-stack system: typed agents, provider abstraction, governed actions — deployed to production.

Frontend

Next.js + TypeScript + Tailwind. Executive Command Center, Workspace, Evidence Intelligence, Decision Ledger, and adaptive experience modes — live on Vercel.

Backend

Python FastAPI + Pydantic schemas, multi-agent orchestrator, deterministic scoring, SQLite decision ledger, HubSpot connector — deployed on Render.

Model & data layer


ModelAdapter provider abstraction (mock today, NVIDIA NIM / Nemotron ready), BYOK session-only keys, and synthetic + HubSpot test-CRM data sources.

Operating model

Amit drives problem framing, architecture, agent contracts, evaluation, and demo story. Saurabh drives app implementation, backend integration, model/API wiring, testing, and deployment readiness.

NVIDIA SDKs, tools, and inference strategy

Provider abstraction ships today; NVIDIA-native runtime plugs into the same typed contract — Current, Hackathon, and Future cleanly separated.



Current (today)

Provider abstraction + BYOK already route to any model. The deterministic engine is the benchmark and fallback; the NVIDIA adapter stub is in place.

Hackathon build

Swap NVIDIA NIM endpoints + Nemotron into the existing adapter for live reasoning and seller drafts — no architecture change required.

Future

NeMo Agent Toolkit orchestration, NemoClaw long-running agents, Triton + GPU optimization, and agent evaluation at scale.

Visible proof in demo

Agent traces, model calls, structured outputs, latency timings, decision ledger, and the human-approval path — all already in the product.

Development and deployment approach

What we did differently: an enterprise-grade workflow around typed agents, observability, and human approval — not a generic chatbot wrapper.

1. Define contracts

Each agent has role, inputs, tools, output schema, allowed data, confidence logic, and escalation conditions.

2. Build dataset service

Synthetic SMB accounts plus a HubSpot test-CRM connector — spend, support, product usage, campaign, and engagement signals with edge cases.

3. Provider abstraction

ModelAdapter routes to any provider; the deterministic engine is the benchmark; NVIDIA NIM / Nemotron plug in through the same typed contract.

4. Add action layer

Rank accounts, explain why, recommend next actions, generate seller-ready call scripts and email drafts, then orchestrate execution.

5. Add governance

Decision ledger with evidence, caveats, confidence, source trace, and approval before action.


6. Deploy and evaluate

Deployed to Vercel + Render with latency profiling, automated evals (10/10), evidence checks, and reviewer-visible agent traces.

Architecture in production: provider-agnostic agent runtime · deployed web app · inference-first · human-in-the-loop by design

Product architecture: governed layers, voice-ready

Voice is an interaction layer above an unchanged governed core — no layer can bypass the one beneath it.



Decision Ledger underneath

An append-only ledger sits beneath every layer, recording evidence, confidence, caveats, approval, and outcome.

Continuous learning loop

Captured business outcomes return as new signals, so reasoning and recommendations improve over time.

Voice never bypasses governance

A spoken request enters the same typed contracts, evidence checks, and human-approval gate as a typed one.

Provider abstraction

A ModelAdapter keeps reasoning swappable — mock today, NVIDIA NIM / Nemotron the moment access lands.

Governance: AI explains, humans decide

Accountability is a feature, not an afterthought — the same governance whether input is typed or spoken.

TRACK

A

Every recommendation is evidence-backed, confidence-scored, and gated by a human. No action reaches a customer without explicit approval.

AI boundary

Agents rank, explain, and draft. They never finalize priority or execute on their own — the model cannot write the decision.

Human approval gate

A seller reviews evidence, confidence, and caveats, then approves, edits, or rejects. Approval is a hard contract, not a suggestion.

Decision Ledger

Append-only single source of truth: query, agents invoked, evidence used, confidence, caveats, approval, and the outcome.

Data & key posture

Synthetic-by-default data; no real customer records. BYOK keys live in session storage only — never persisted, logged, returned, or deployed — with a deterministic fallback always available.

Revenue Execution Center

The execution layer that turns an approved decision into a measured business result.

TRACK

A

1

Execution Agent

Orchestrates the approved action — outreach, scheduling, follow-up — as a guided, event-driven flow.

2

Customer Response

The account responds: meeting accepted, reply received, or no-response captured.

3

Outcome

The result is recorded against the recommendation — win, partial, or escalation needed.

4

Ledger Updated

Execution events stream into the Decision Ledger, Lifecycle ribbon, and Timeline automatically.

5

Business Outcome

Stated in business terms: meeting booked, renewal risk reduced, opportunity created, customer engaged.

Business outcomes, not workflow steps

Auto-advance or manual control — the seller watches an orchestrated flow, not a wizard to push through.

Every execution event is governed and logged; nothing writes back to the live CRM in demo mode.

Closes the operating loop: signal → recommendation → approval → execution → outcome → ledger.

The Voice Chief of Staff (planned hackathon)

Voice becomes another enterprise interaction channel — not another AI model. The governed core stays unchanged.



Planned hackathon implementation

Gnani.ai is our planned speech layer for the hackathon build — not yet implemented. The architecture is voice-ready today.

Enterprise speech

Low-latency, telephony-optimized recognition built for noisy, real-world enterprise conditions.

Indian languages + code-switching

Multilingual support with English / regional code-switching for Indian enterprise revenue teams.

Product evolution

Revenue OS → Enterprise AI Chief of Staff → voice-native Revenue OS. Voice is an interaction layer, not a new brain.

Business value: the executive case

Every capability maps to a number a revenue leader already owns — this reads like a CIO business case.

Revenue at risk

Surface at-risk renewals and declining accounts early, while there is still time to intervene.

Revenue saved

Recovery plans on flagged accounts protect renewals that would otherwise churn silently.

Pipeline protected

Expansion and cross-sell signals are caught and acted on, not missed in a crowded book.

Seller productivity

Minutes, not days: brief, evidence, and draft are prepared so sellers act instead of stitching context.

Decision cycle time

Signal to approved action in one governed loop, compressing the time from insight to outcome.

Governed & explainable

Every decision is evidence-backed, confidence-scored, human-approved, and auditable in the ledger.

Outcome-measured by design: revenue protected, pipeline created, and decision cycle time — all attributable in the Decision Ledger.

Potential bottlenecks and mitigation plan

The main risk is not building “more agents”; it is building a reliable, explainable workflow that stays fast and reviewable.

Bottleneck 1: NVIDIA access path

Mitigation: provider abstraction already runs on a deterministic engine; NVIDIA NIM / Nemotron plug in via the same adapter when access lands — zero rework.

Bottleneck 2: orchestration noise

Mitigation: bounded agents, typed outputs, fixed workflow graph, validation checks, and no open-ended autonomous execution.

Bottleneck 3: latency

Mitigation: parallelize independent agents, cache summaries, reduce unnecessary model calls, and profile each step.

Bottleneck 4: data realism

Mitigation: synthetic accounts with missing fields, conflicting signals, stale notes, noisy engagement, and realistic prioritization conflicts.

Bottleneck 5: trust quality

Mitigation: evidence-first recommendation format, confidence scoring, caveats, human approval, and rejection/edit capture.

Bottleneck 6: scope control

Mitigation: one sharp workflow: rank accounts → explain why → recommend action → draft message → approval ledger.



Road Map

Ideas for further improvement of the solution.



Roadmap: three horizons + execution plan

Where the product is today, what we build at the hackathon, and where it goes next — one continuous arc.

Where we want NVIDIA / OpenACC mentor support

Current — shipped today

Executive Command Center, Portfolio Intelligence, Revenue Execution Center, Governance, and Decision Ledger.

Hackathon — voice

Voice Chief of Staff with Gnani.ai: voice conversations, multilingual support, and spoken portfolio reviews over the governed core.

Future — vision

Digital executive avatar, live meeting coach, and an enterprise multimodal workspace.

NVIDIA acceleration

NIM + Nemotron reasoning, NeMo orchestration, and Triton + GPU optimization plug into every horizon through the provider layer.

Execution roadmap



Final demo target: “Which accounts need attention this week?” → ranked actions with evidence, confidence, voice summary, approval controls, and visible NVIDIA model/runtime calls.

Mentor ask: where NVIDIA guidance accelerates us

A deployed product today, with a credible path to a reusable VentureOS action-layer workflow later.

Where we want NVIDIA mentorship

1. NeMo + Nemotron

NeMo Agent Toolkit orchestration patterns and Nemotron tuning for reasoning, ranking, and structured action generation.

2. NIM + Triton deployment

NIM endpoint deployment and Triton optimization for low-latency, GPU-backed inference across the agent workflow.

3. Voice pipeline + SALM

Speech pipeline optimization and SALM integration guidance for the planned Gnani.ai voice layer — low-latency, multilingual, enterprise voice.

Key limitations and how we will improve them

Limitation	Improvement plan
Synthetic data only	Synthetic-by-default plus a HubSpot test CRM today; connect permitted enterprise sources later with privacy and access controls.
Provider abstraction	Deterministic engine is the benchmark; NVIDIA NIM / Nemotron and BYOK providers plug in through the same typed contract.
Agent reliability risk	Use typed outputs, bounded steps, explicit fallback logic, confidence thresholds, and human approval before any action.
Latency and cost	Parallelize independent agents, cache account summaries, profile NIM calls, and use mentor support for GPU-backed optimization.

Long-term ambition: make agentic AI accountable enough for enterprise action — accelerated by NVIDIA, governed by design, and approved by humans.

Closing summary

From systems of record to systems of reasoning.

Clear real-world use case

Enterprise teams struggle to convert fragmented signals into timely, explainable action. Signal-to-Action is deployed and demo-ready today.

Strong Track A fit

The solution is a controlled multi-agent workflow, not a generic chatbot. It has orchestration, specialized agents, tool routing, traces, and approval.

NVIDIA-centered build

Provider abstraction makes NVIDIA NIM / Nemotron a drop-in runtime; optimization and evaluation are first-class goals.

We believe enterprise software is evolving from systems of record to systems of reasoning. Signal-to-Action Agent is our first step toward an Enterprise AI Chief of Staff — a governed, explainable, voice-native operating system that helps every revenue team move confidently from signal to business outcome.

Signal → Reasoning → Approval → Execution → Outcome



Need Help? Contact Us!!

Join the Slack Workspace [here](#)

